

# Infrastructure Provisioning with AWS & CI CD Objectives

## Introduction

This guide outlines the steps to provision a basic AWS infrastructure for hosting a web application. The setup includes creating an EC2 instance, configuring security settings to allow HTTP traffic, and establishing a Virtual Private Cloud (VPC) with a public subnet. Additionally, the guide details the integration of a CI/CD pipeline using GitLab to automate the deployment of the application. Automation options using Terraform are also included for efficient and repeatable deployments.

### 1. Overview

This guide provides instructions to:

- **Set up AWS infrastructure** using Terraform.
  - **Build, push, and deploy a Flask application** containerized with Docker to AWS ECS.
  - **Test endpoints and validate the deployment.**
- 

### 2. AWS Infrastructure Setup

#### Prerequisites

- **AWS CLI** installed and configured with sufficient permissions.
- **Terraform CLI** installed.

#### Steps

##### Initialize Terraform

```
terraform init
```

##### Review and Apply Terraform Configurations

- Verify the resources defined in the `.tf` files.

Apply the configurations:

```
terraform plan
```

```
Terraform apply
```

- This will create:
  - VPC, subnets, and internet gateway (`network.tf`).
  - ECS task execution IAM role and policy (`ecs-iam.tf` and `iam.tf`).
  - An ECR repository for the application (`ecr.tf`).
  - Outputs the ECR repository URL (`outputs.tf`).

### Verify Outputs

- Copy the `ecr_repository_url` from the output. You'll use this URL to push your Docker image.
- 

## 3. Flask Application

### Application Details

- The Flask application provides:
  - A health check endpoint: `/health`
  - A root endpoint: `/`

### Code Structure

- `app.py`: Main Flask app code with logging and endpoints.
  - `test_app.py`: Pytest-based tests for the Flask application.
- 

## 4. Containerizing the Flask App

### Steps

**Build the Docker Image** Navigate to the `python-app` directory and build the image:

bash

Copy code

```
cd python-app
```

```
docker build -t flask-app .
```

---

## 5. Deploying to AWS ECS

### Infrastructure Assumptions

- The ECS cluster, ALB, and related services are created via Terraform (defined in `alb.tf` and other files).

## Steps

### 1. Create an ECS Task Definition

- Use the `ecs-task-execution-role` for execution.
- Reference the ECR image `<ECR_REPOSITORY_URL>:latest`.

### 2. Deploy the ECS Service

- Attach the service to the ALB.
- Ensure the ALB forwards traffic from port `80` to your ECS tasks (on port `5000`).

---

## 6. Testing the Deployment

### Verify Application Endpoints

1. Fetch the ALB's DNS name or public IP (from the AWS console or Terraform output).
2. Test the endpoints using `curl` or a browser:

Root endpoint:

```
curl http://<ALB-DNS-NAME>/
```

Expected output:

```
json
{"message": "Hello, World!"}
```

◦

Health check endpoint:

```
curl http://<ALB-DNS-NAME>/health
```

Expected output:

```
json
{
  "status": "healthy",
  "timestamp": "...",
  "version": "1.0.0",
  "python_version": "...",
  "environment": "production"
}
```

### Run Tests Locally

Navigate to `python-app` and run the tests:

```
pytest app/test_app.py
```

Expected output:

```
2 passed in 0.XXs
```

---

## 7. CI/CD Pipeline

Define CI/CD Pipeline in `ci-cd.yaml`

- Automate the following tasks:
  1. Run unit tests.
  2. Build the Docker image.
  3. Push the image to ECR.
  4. Deploy updated tasks to ECS.

## 8. Key Notes

- **Security Best Practices:**
  - Restrict IAM roles and security group rules to the least privilege.
  - Use HTTPS for ALB listeners with an ACM certificate.
- **Scalability:**
  - Update ECS service configurations for auto-scaling based on CPU/memory usage.
- **Monitoring:**
  - Enable CloudWatch logging for ECS tasks.
  - Set up ALB access logs for detailed insights.