

# Kubernetes Automatic Scaling and Monitoring Documentation

## Table of Contents

- Overview
- Prerequisites
- Installation
- Configuration
- Usage Guide
- Monitoring Features
- Scaling Logic
- Logging and Metrics
- Troubleshooting
- Best Practices

## 1. Overview

This automation tool provides continuous monitoring and automatic scaling of Kubernetes deployments based on resource utilization. It helps maintain optimal performance by automatically adjusting the number of replicas based on CPU and memory usage.

### Key Features

- Automatic resource monitoring
- Dynamic scaling based on configurable thresholds
- Real-time metrics dashboard
- Historical metrics logging
- Color-coded status outputs
- Configurable scaling parameters

## 2. Prerequisites

### Required components:

—>Kubernetes CLI

**kubectrl version --client**

->Basic Calculator (for arithmetic operations)

**bc -v**

### Required Kubernetes permissions

- kubectl auth can-i scale deployment
- kubectl auth can-i get pods
- kubectl auth can-i get deployments

## 3. Installation

### 1. Download the script:

```
curl -O https://raw.githubusercontent.com/your-repo/k8s-auto-scale.sh
```

### 2. Make it executable:

```
chmod +x k8s-auto-scale.sh
```

### 3. Verify installation:

```
./k8s-auto-scale.sh --version
```

## 4. Configuration

->Default Parameters

```
CPU_THRESHOLD=80      # CPU usage percentage threshold
MEMORY_THRESHOLD=80   # Memory usage percentage threshold
MAX_REPLICAS=10       # Maximum number of replicas
MIN_REPLICAS=2        # Minimum number of replicas
CHECK_INTERVAL=30     # Monitoring interval in seconds
```

### Custom Configuration

You can override defaults during script execution:

```
Enter CPU threshold % (default 80): 70
Enter Memory threshold % (default 80): 75
Enter minimum replicas (default 2): 3
Enter maximum replicas (default 10): 8
Enter check interval in seconds (default 30): 45
```

## 5. Usage Guide

->Basic Usage

```
./k8s-auto-scale.sh
```

## Required Inputs

Enter namespace: production

Enter deployment name: web-app

## 6. Monitoring Features

### Real-time Metrics

- CPU usage percentage
- Memory usage percentage
- Current replica count
- Resource utilization trends

### Dashboard Display

#### Current Metrics Summary:

-----

Last 5 measurements:

2024-11-18 10:30:00 CPU: 45% Memory: 60% Replicas: 3

2024-11-18 10:30:30 CPU: 48% Memory: 62% Replicas: 3

2024-11-18 10:31:00 CPU: 82% Memory: 65% Replicas: 4

2024-11-18 10:31:30 CPU: 75% Memory: 63% Replicas: 4

2024-11-18 10:32:00 CPU: 70% Memory: 61% Replicas: 4

-----

...

## 7. Scaling Logic

### Scale Up Conditions

if (CPU\_Usage > CPU\_THRESHOLD) OR (Memory\_Usage > MEMORY\_THRESHOLD):

    if (Current\_Replicas < MAX\_REPLICAS):

        Scale Up by 1 replica

### Scale Down Conditions

if (CPU\_Usage < CPU\_THRESHOLD/2) AND (Memory\_Usage < MEMORY\_THRESHOLD/2):

    if (Current\_Replicas > MIN\_REPLICAS):

        Scale Down by 1 replica

## 8. Logging and Metrics

### Log File Format

`metrics_${namespace}_${deployment}.log`

### Sample Log Entry

`2024-11-18 10:30:00 CPU: 45% Memory: 60% Replicas: 3`

### Metrics Collected

- Timestamp
- CPU usage percentage
- Memory usage percentage
- Number of replicas
- Scaling events

## 9. Troubleshooting

### Common Issues and Solutions

#### 1. Permission Issues

Error: "cannot get deployments"

Solution: Ensure proper RBAC permissions

#### 2. Resource Metrics Unavailable

Error: "metrics not available"

Solution: Verify metrics-server is running

#### 3. Scaling Failures

Error: "scaling deployment failed"

Solution: Check deployment configuration and resource quotas

## 10. Best Practices

### Resource Thresholds

- Set CPU threshold based on application behavior
- Consider memory usage patterns
- Adjust thresholds during peak hours

### Scaling Configuration

- Set appropriate MIN\_REPLICAS for HA
- Configure MAX\_REPLICAS based on cluster capacity

- Use reasonable CHECK\_INTERVAL (30-60 seconds)

### Monitoring

- Regular log review
- Track scaling patterns
- Monitor application performance

### Maintenance

- Regular script updates
- Log rotation
- Performance tuning

### Health Checks

The script performs the following health checks:

#### 1. Deployment Health

```
kubectl rollout status deployment/$deployment -n $namespace
```

#### 2. Resource Availability

```
kubectl top pods -n $namespace
```

#### 3. Scaling Operations

```
kubectl scale deployment $deployment --replicas=$count -n $namespace
```

### Security Considerations

#### 1. RBAC Permissions

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: Role
```

```
metadata:
```

```
  name: deployment-scaler
```

```
rules:
```

```
- apiGroups: ["apps"]
```

```
  resources: ["deployments"]
```

```
  verbs: ["get", "list", "update", "patch"]
```

```
- apiGroups: [""]
```

```
  resources: ["pods"]
```

```
  verbs: ["get", "list"]
```

#### 2. Namespace Isolation

- Use namespace-scoped roles
- Limit access to specific deployments

- Monitor scaling operations

## **Limitations**

### **1. Resource Constraints**

- Dependent on metrics-server availability
- May have delayed scaling during heavy load
- Limited by cluster resources

### **2. Scaling Boundaries**

- Minimum replica count enforced
- Maximum replica count enforced
- Scaling step size fixed at 1

## **Support and Maintenance**

### **Updates and Patches**

- Regular script updates
- Bug fixes
- Feature enhancements

### **Monitoring and Alerts**

- Resource usage alerts
- Scaling event notifications
- Error reporting

This documentation provides a detail guide for using the automatic scaling script.